

[Click to Take the FREE Imbalanced Classification Crash-Course](#)



Cost-Sensitive Logistic Regression for Imbalanced Classification

by Jason Brownlee on [January 27, 2020](#) in [Imbalanced Classification](#)

[Tweet](#)[Share](#)[Share](#)

Last Updated on February 1, 2020

Logistic regression does not support imbalanced classification directly.

Instead, the training algorithm used to fit the logistic regression model must be modified to take the skewed distribution into account. This can be achieved by specifying a class weighting configuration that is used to influence the amount that logistic regression coefficients are updated during training.

The weighting can penalize the model less for errors made on examples from the majority class and penalize the model more for errors made on examples from the minority class. The result is a version of logistic regression that performs better on imbalanced classification tasks, generally referred to as cost-sensitive or weighted logistic regression.

In this tutorial, you will discover cost-sensitive logistic regression for imbalanced classification.

After completing this tutorial, you will know:

- How standard logistic regression does not support imbalanced classification.
- How logistic regression can be modified to weight the classes using class weights or coefficients.
- How to configure class weight for logistic regression and how to choose appropriate configurations.

Discover SMOTE, one-class classification, cost-sensitive learning, threshold shifting, and much more in my new book, with 30 step-by-step tutorials and full Python code.

Let's get started.

Start Machine Learning

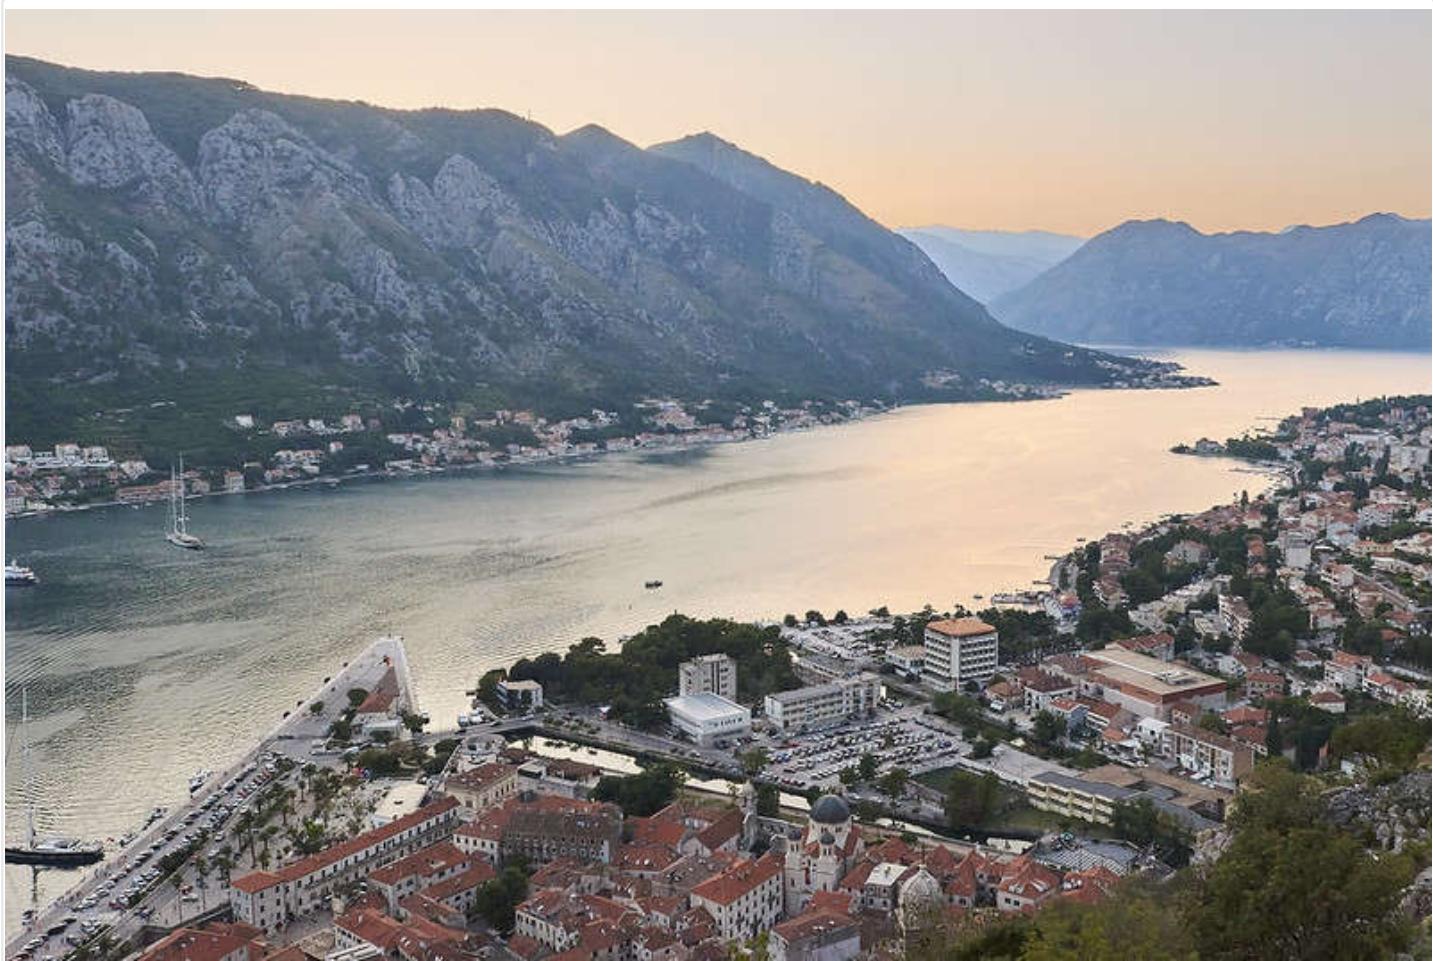
You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

my

- Update Feb/2020: Fixed typo in weight calculation.



Cost-Sensitive Logistic Regression for Imbalanced Classification

Photo by [Naval S](#), some rights reserved.

Tutorial Overview

This tutorial is divided into five parts; they are:

1. Imbalanced Classification Dataset
2. Logistic Regression for Imbalanced Classification
3. Weighted Logistic Regression With Scikit-Learn
4. Grid Search Weighted Logistic Regression

Imbalanced Classification Data

Before we dive into the modification of logistic regression for imbalanced classification dataset.

We can use the `make_classification()` function to define the dataset. We will generate 10,000 examples with an ap-

```
1 ...  
2 # define dataset
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

```
3 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
4     n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
```

Once generated, we can summarize the class distribution to confirm that the dataset was created as we expected.

```
1 ...
2 # summarize class distribution
3 counter = Counter(y)
4 print(counter)
```

Finally, we can create a scatter plot of the examples and color them by class label to help understand the challenge of classifying examples from this dataset.

```
1 ...
2 # scatter plot of examples by class label
3 for label, _ in counter.items():
4     row_ix = where(y == label)[0]
5     pyplot.scatter(X[row_ix, 0], X[row_ix, 1], label=str(label))
6 pyplot.legend()
7 pyplot.show()
```

Tying this together, the complete example of generating the synthetic dataset and plotting the examples is listed below.

```
1 # Generate and plot a synthetic imbalanced classification dataset
2 from collections import Counter
3 from sklearn.datasets import make_classification
4 from matplotlib import pyplot
5 from numpy import where
6 # define dataset
7 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
8     n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
9 # summarize class distribution
10 counter = Counter(y)
11 print(counter)
12 # scatter plot of examples by class label
13 for label, _ in counter.items():
14     row_ix = where(y == label)[0]
15     pyplot.scatter(X[row_ix, 0], X[row_ix, 1], label=str(label))
16 pyplot.legend()
17 pyplot.show()
```

Running the example first creates the dataset and summarizes the class distribution.

We can see that the dataset has an approximate 1:10 ratio of examples in the majority class and 100 in the minority

```
1 Counter({0: 9900, 1: 100})
```

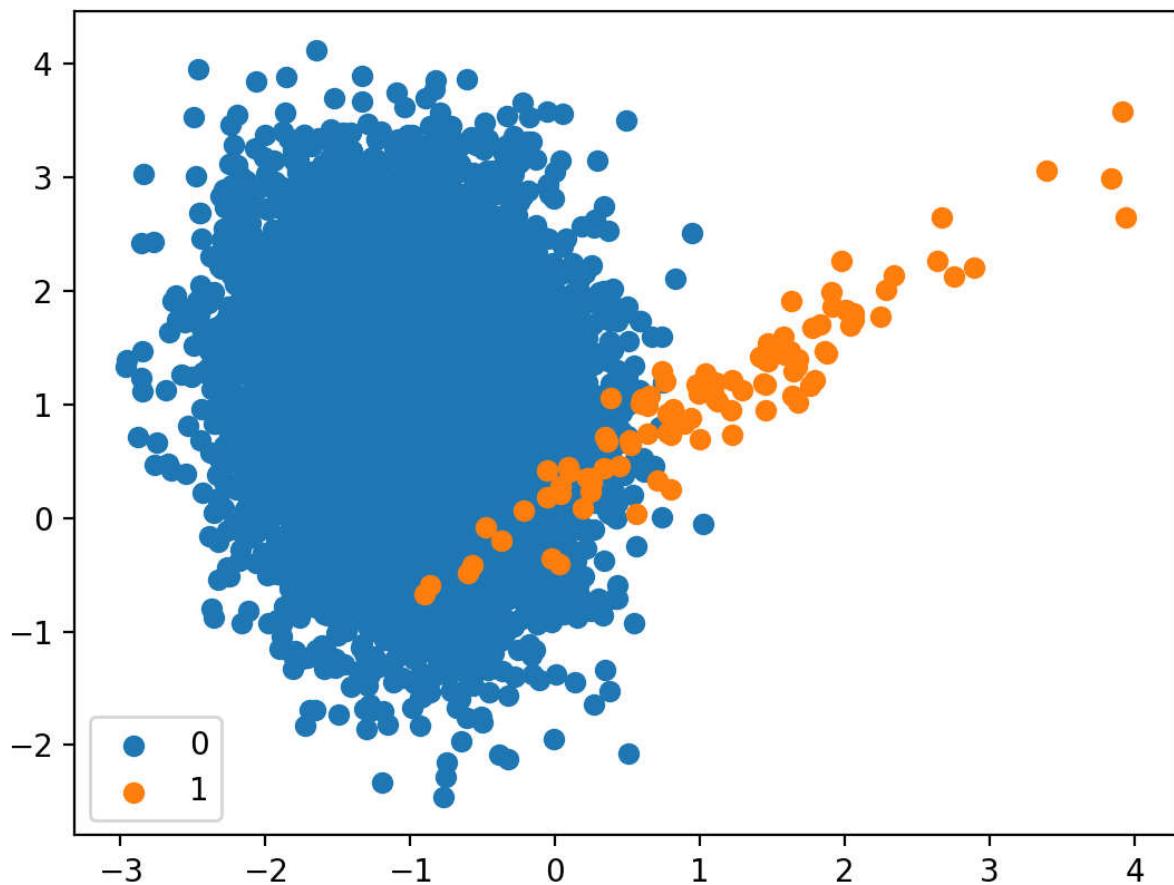
Next, a scatter plot of the dataset is created showing the large number of examples for the majority class (blue) and a small number of examples for the minority class (orange).

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)



Scatter Plot of Binary Classification Dataset With 1 to 100 Class Imbalance

Next, we can fit a standard logistic regression model on the dataset.

We will use repeated cross-validation to evaluate the model, with three repeats of 10-fold cross-validation. The mode performance will be reported using the mean ROC area under curve (ROC AUC) averaged over repeats and all folds.

```

1 ...
2 # define evaluation procedure
3 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
4 # evaluate model
5 scores = cross_val_score(model, X, y, scoring='roc_auc')
6 # summarize performance
7 print('Mean ROC AUC: %.3f' % mean(scores))

```

Tying this together, the complete example of evaluating a logistic regression model on an imbalanced classification problem is listed below.

```

1 # fit a logistic regression model on an imbalanced dataset
2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import cross_val_score
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from sklearn.linear_model import LogisticRegression

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```

7 # generate dataset
8 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
9   n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
10 # define model
11 model = LogisticRegression(solver='lbfgs')
12 # define evaluation procedure
13 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
14 # evaluate model
15 scores = cross_val_score(model, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
16 # summarize performance
17 print('Mean ROC AUC: %.3f' % mean(scores))

```

Running the example evaluates the standard logistic regression model on the imbalanced dataset and reports the mean ROC AUC.

We can see that the model has skill, achieving a ROC AUC above 0.5, in this case achieving a mean score of 0.985.

1 Mean ROC AUC: 0.985

This provides a baseline for comparison for any modifications performed to the standard logistic regression algorithm.

Want to Get Started With Imbalance Classification?

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

[Download Your FREE Mini-Course](#)

Logistic Regression for Imbalanced Classification

Logistic regression is an effective model for binary classification tasks, although by default, it is not effective at imbalanced classification.

Logistic regression can be modified to be better suited

The coefficients of the logistic regression algorithm are updated by minimizing the negative log likelihood (loss) for the model on the training data.

- minimize sum i to n -(log(yhat_i) * y_i + log(1 - yhat_i) * (1 - y_i))

This involves the repeated use of the model to make predictions, then update the coefficients in a direction that reduces the loss of the model.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

[START MY EMAIL COURSE](#)

The calculation of the loss for a given set of coefficients can be modified to take the class balance into account.

By default, the errors for each class may be considered to have the same weighting, say 1.0. These weightings can be adjusted based on the importance of each class.

- minimize sum i to n -(w0 * log(yhat_i) * y_i + w1 * log(1 - yhat_i) * (1 - y_i))

The weighting is applied to the loss so that smaller weight values result in a smaller error value, and in turn, less update to the model coefficients. A larger weight value results in a larger error calculation, and in turn, more update to the model coefficients.

- **Small Weight:** Less importance, less update to the model coefficients.
- **Large Weight:** More importance, more update to the model coefficients.

As such, the modified version of logistic regression is referred to as Weighted Logistic Regression, Class-Weighted Logistic Regression or Cost-Sensitive Logistic Regression.

The weightings are sometimes referred to as importance weightings.

Although straightforward to implement, the challenge of weighted logistic regression is the choice of the weighting to use for each class.

Weighted Logistic Regression with Scikit-Learn

The scikit-learn Python machine learning library provides an implementation of logistic regression that supports class weighting.

The [LogisticRegression class](#) provides the `class_weight` argument that can be specified as a model hyperparameter. The `class_weight` is a dictionary that defines each class label (e.g. 0 and 1) and the weighting to apply in the calculation of the negative log likelihood when fitting the model.

For example, a 1 to 1 weighting for each class 0 and 1 can be defined as follows:

```
1 ...
2 # define model
3 weights = {0:1.0, 1:1.0}
4 model = LogisticRegression(solver='lbfgs', cla
```

The class weighing can be defined multiple ways; for example:

- **Domain expertise**, determined by talking to subject matter experts.
- **Tuning**, determined by a hyperparameter search.
- **Heuristic**, specified using a general best practice.

A best practice for using the class weighting is to use domain expertise to determine the class weightings for the training dataset.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

For example, the class distribution of the test dataset is a 1:100 ratio for the minority class to the majority class. The inversion of this ratio could be used with 1 for the majority class and 100 for the minority class; for example:

```
1 ...
2 # define model
3 weights = {0:1.0, 1:100.0}
4 model = LogisticRegression(solver='lbfgs', class_weight=weights)
```

We might also define the same ratio using fractions and achieve the same result; for example:

```
1 ...
2 # define model
3 weights = {0:0.01, 1:1.0}
4 model = LogisticRegression(solver='lbfgs', class_weight=weights)
```

We can evaluate the logistic regression algorithm with a class weighting using the same evaluation procedure defined in the previous section.

We would expect that the class-weighted version of logistic regression to perform better than the standard version of logistic regression without any class weighting.

The complete example is listed below.

```
1 # weighted logistic regression model on an imbalanced classification dataset
2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import cross_val_score
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from sklearn.linear_model import LogisticRegression
7 # generate dataset
8 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
9     n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
10 # define model
11 weights = {0:0.01, 1:1.0}
12 model = LogisticRegression(solver='lbfgs', class_weight=weights)
13 # define evaluation procedure
14 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
15 # evaluate model
16 scores = cross_val_score(model, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
17 # summarize performance
18 print('Mean ROC AUC: %.3f' % mean(scores))
```

Running the example prepares the synthetic imbalanced classification dataset, then evaluates the class-weighted version of logistic regression using repeated

The mean ROC AUC score is reported, in this case showing the class-weighted logistic regression, 0.989 as compared to 0.985.

```
1 Mean ROC AUC: 0.989
```

The scikit-learn library provides an implementation of this class weighting.

It is implemented via the `compute_class_weight()` function.

- `n_samples / (n_classes * n_samples_with_class)`

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

We can test this calculation manually on our dataset. For example, we have 10,000 examples in the dataset, 9900 in class 0, and 100 in class 1.

The weighting for class 0 is calculated as:

- weighting = $n_samples / (n_classes * n_samples_with_class)$
- weighting = $10000 / (2 * 9900)$
- weighting = $10000 / 19800$
- weighting = 0.05

The weighting for class 1 is calculated as:

- weighting = $n_samples / (n_classes * n_samples_with_class)$
- weighting = $10000 / (2 * 100)$
- weighting = $10000 / 200$
- weighting = 50

We can confirm these calculations by calling the `compute_class_weight()` function and specifying the `class_weight` as “*balanced*.” For example:

```
1 # calculate heuristic class weighting
2 from sklearn.utils.class_weight import compute_class_weight
3 from sklearn.datasets import make_classification
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
6   n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
7 # calculate class weighting
8 weighting = compute_class_weight('balanced', [0,1], y)
9 print(weighting)
```

Running the example, we can see that we can achieve a weighting of about 0.5 for class 0 and a weighting of 50 for class 1.

These values match our manual calculation.

```
1 [ 0.50505051 50. ]
```

The values also match our heuristic calculation above for inverting the ratio of the class distribution in the training dataset; for example:

- $0.5:50 == 1:100$

We can use the default class balance directly with the `class_weight` argument to ‘*balanced*.’ For example:

```
1 ...
2 # define model
3 model = LogisticRegression(solver='lbfgs', class_weight='balanced')
```

The complete example is listed below.

```
1 # weighted logistic regression for class imbalance
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```

2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import cross_val_score
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from sklearn.linear_model import LogisticRegression
7 # generate dataset
8 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
9     n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
10 # define model
11 model = LogisticRegression(solver='lbfgs', class_weight='balanced')
12 # define evaluation procedure
13 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
14 # evaluate model
15 scores = cross_val_score(model, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
16 # summarize performance
17 print('Mean ROC AUC: %.3f' % mean(scores))

```

Running the example gives the same mean ROC AUC as we achieved by specifying the inverse class ratio manually.

```
1 Mean ROC AUC: 0.989
```

Grid Search Weighted Logistic Regression

Using a class weighting that is the inverse ratio of the training data is just a heuristic.

It is possible that better performance can be achieved with a different class weighting, and this too will depend on the choice of performance metric used to evaluate the model.

In this section, we will grid search a range of different class weightings for weighted logistic regression and discover which results in the best ROC AUC score.

We will try the following weightings for class 0 and 1:

- {0:100,1:1}
- {0:10,1:1}
- {0:1,1:1}
- {0:1,1:10}
- {0:1,1:100}

These can be defined as grid search parameters for the model.

```

...
# define grid
balance = [{0:100,1:1}, {0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}]
param_grid = dict(class_weight=balance)

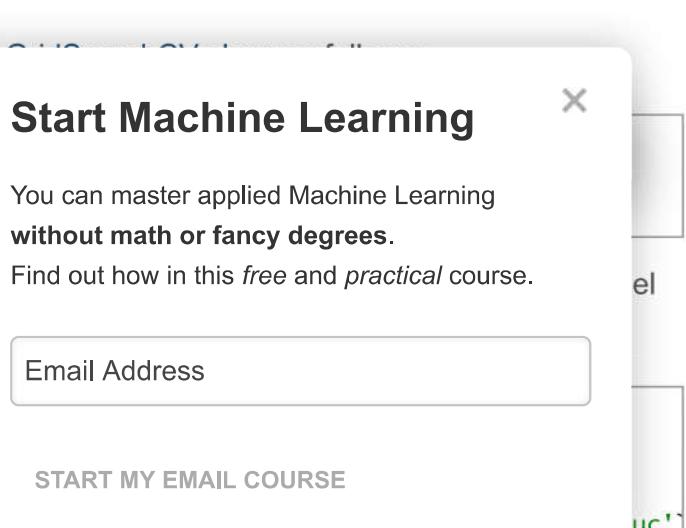
```

We can perform the grid search on these parameters and evaluate the performance using ROC AUC:

```

1 ...
2 # define evaluation procedure
3 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
4 # define grid search
5 grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring='roc_auc', cv=cv, n_jobs=-1)

```



Once executed, we can summarize the best configuration as well as all of the results as follows:

```
1 ...
2 # report the best configuration
3 print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
4 # report all configurations
5 means = grid_result.cv_results_['mean_test_score']
6 stds = grid_result.cv_results_['std_test_score']
7 params = grid_result.cv_results_['params']
8 for mean, stdev, param in zip(means, stds, params):
9     print("%f (%f) with: %r" % (mean, stdev, param))
```

Tying this together, the example below grid searches five different class weights for logistic regression on the imbalanced dataset.

We might expect that the heuristic class weighing is the best performing configuration.

```
1 # grid search class weights with logistic regression for imbalance classification
2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import GridSearchCV
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from sklearn.linear_model import LogisticRegression
7 # generate dataset
8 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
9     n_clusters_per_class=1, weights=[0.99], flip_y=0, random_state=2)
10 # define model
11 model = LogisticRegression(solver='lbfgs')
12 # define grid
13 balance = [{0:100,1:1}, {0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}]
14 param_grid = dict(class_weight=balance)
15 # define evaluation procedure
16 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
17 # define grid search
18 grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=cv, scoring='roc_auc')
19 # execute the grid search
20 grid_result = grid.fit(X, y)
21 # report the best configuration
22 print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
23 # report all configurations
24 means = grid_result.cv_results_['mean_test_score']
25 stds = grid_result.cv_results_['std_test_score']
26 params = grid_result.cv_results_['params']
27 for mean, stdev, param in zip(means, stds, params):
28     print("%f (%f) with: %r" % (mean, stdev, param))
```

Running the example evaluates each class weighting, and prints the best configuration and the associated mean ROC AUC score.

In this case, we can see that the 1:100 majority to minority class weighting resulted in the highest mean ROC AUC score. This matches the configuration for the general 1:100 imbalance.

It might be interesting to explore even more severe class imbalances to see how the ROC AUC score changes.

```
1 Best: 0.989077 using {'class_weight': {0: 1, 1: 100}}
2 0.982498 (0.016722) with: {'class_weight': {0: 1, 1: 100}}
3 0.983623 (0.015760) with: {'class_weight': {0: 1, 1: 100}}
4 0.985387 (0.013890) with: {'class_weight': {0: 1, 1: 100}}
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```
5 0.988044 (0.010384) with: {'class_weight': {0: 1, 1: 10}}
6 0.989077 (0.006865) with: {'class_weight': {0: 1, 1: 100}}
```

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Papers

- Logistic Regression in Rare Events Data, 2001.
- The Estimation of Choice Probabilities from Choice Based Samples, 1977.

Books

- Learning from Imbalanced Data Sets, 2018.
- Imbalanced Learning: Foundations, Algorithms, and Applications, 2013.

APIs

- `sklearn.utils.class_weight.compute_class_weight` API.
- `sklearn.linear_model.LogisticRegression` API.
- `sklearn.model_selection.GridSearchCV` API.

Summary

In this tutorial, you discovered cost-sensitive logistic regression for imbalanced classification.

Specifically, you learned:

- How standard logistic regression does not support imbalanced classification.
- How logistic regression can be modified to weight model error by class weight when fitting the coefficients.
- How to configure class weight for logistic regression and how to grid search different class weight configurations.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer them.

Get a Handle on Imbalanced Classification

Develop Imbalanced Learning Models

...with just a few lines of code

Discover how to handle
Imbalanced Classification

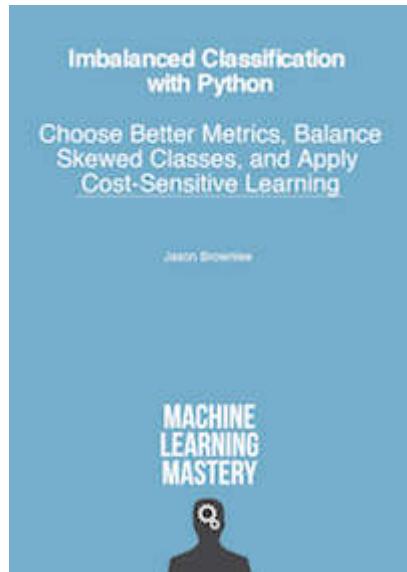
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



It provides **self-study tutorials** and **end-to-end projects** on:
Performance Metrics, Undersampling Methods, SMOTE, Threshold Moving, Probability Calibration, Cost-Sensitive Algorithms
and much more...

Bring Imbalanced Classification Methods to Your Machine Learning Projects

SEE WHAT'S INSIDE

Tweet

Share

Share



About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

◀ [Tour of Data Sampling Methods for Imbalanced Classification](#)

[Cost-Sensitive Decision Trees for Imbalanced Classification](#) ▶

18 Responses to *Cost-Sensitive Logistic Regression for Imbalanced Classification*



Elie January 27, 2020 at 7:11 am #

Jason, almost done from reading the book.

Really great piece of work!

One minor recommendation: I'd like to see more expla



Jason Brownlee January 27, 2020 at 7:40 am

Well done on your progress!

REPLY ↗

Start Machine Learning

You can master applied Machine Learning

without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Great suggestion, thanks.



Rodney Silva January 27, 2020 at 10:18 am #

REPLY ↗

I would like to see the improvement on the precision and recall for the minor class.



Jason Brownlee January 27, 2020 at 2:33 pm #

REPLY ↗

Thanks, great suggestion.



marco January 29, 2020 at 2:42 am #

REPLY ↗

Hello Jason,

a question about SVC and linearSVC.

What is the difference?

I'm trying a sentiment analysis with 1000 observations (750 training + 250 test).

Is it better to use SVC or linearSVC for the analysis?

What is the meaning of C hyperparameter in linearSVC (in simple words)?

I've found the C parameter is common also in other algorithms. The meaning is the same?

Thanks



Jason Brownlee January 29, 2020 at 6:45 am #

REPLY ↗

SVC can do linear SVC via a linear kernel. Linear SVC just uses the linear kernel and nothing else and is optimized for this use case – faster/more efficient.

More on C here:

<https://machinelearningmastery.com/support-vector-machines-for-machine-learning/>



Diane Halliwell January 30, 2020 at 6:58 am #

Hi

Does your book cover example-dependent cost-sensitivity?



Jason Brownlee January 30, 2020 at 6:59 am #

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

No, just classes-based costs.



Temitope Mamukuyomi January 31, 2020 at 8:22 am #

REPLY ↗

Thanks for the tutorial



Jason Brownlee January 31, 2020 at 2:04 pm #

REPLY ↗

You're very welcome, I hope it helps you with your project!



macilane manjate January 31, 2020 at 3:43 pm #

REPLY ↗

Dear Jason Brownlee,

Good morning.

Whren will you have the book in R.

Kind regards,

Macilane



Jason Brownlee February 1, 2020 at 5:46 am #

REPLY ↗

No plans at this stage. My focus is Python given that it is the most popular language for machine learning at the moment.



Sergio Garcia Garcia January 31, 2020 at 7:53 pm #

REPLY ↗

"For example, we have 10,000 examples in the dataset, 9990 in class 0, and 100 in class 1."

Wouldn't be 9990-10 or 9900-100?

Good article, clear explanation



Jason Brownlee February 1, 2020 at 5:53 am #

Thanks, fixed!

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Carlos February 7, 2020 at 3:20 am #

REPLY ↗

Could you please provide the code for "Counter"

NameError: name 'Counter' is not defined

Thank you



Jason Brownlee February 7, 2020 at 8:24 am #

REPLY ↗

You must copy the full code example that includes the important statement.



Atefeh April 12, 2020 at 4:19 am #

REPLY ↗

Thank you for the great work. It was very helpful. I was looking to see how we can approach a specific problem both with classification and regression models? Let's say I want to study the air temperature-soil temperature relation both with classification and regression models to get both the best fit and the decision boundary. Is there any article that can illustrate how steps will be different for these two approaches?



Jason Brownlee April 12, 2020 at 6:25 am #

REPLY ↗

You're welcome.

I don't have an example, sorry.

Leave a Reply

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Name (required)

Email (will not be published) (required) Website[SUBMIT COMMENT](#)

Welcome!

My name is *Jason Brownlee PhD*, and I **help developers** get results with **machine learning**.

[Read more](#)

Never miss a tutorial:



Picked for you:

[8 Tactics to Combat Imbalanced Classes in Your Machine Learning Dataset](#)[Imbalanced Classification With Python \(7-Day Mini-Course\)](#)[SMOTE for Imbalanced Classification with Python](#)[Tour of Evaluation Metrics for Imbalanced Classification](#)[Imbalanced Multiclass Classification with the Glass](#)

Loving the

[The Imbalanced Classification with Python E](#)

Start Machine Learning

[X](#)

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address[START MY EMAIL COURSE](#)

[SEE WHAT'S INSIDE](#)

© 2019 Machine Learning Mastery Pty. Ltd. All Rights Reserved.
Address: PO Box 206, Vermont Victoria 3133, Australia. | ACN: 626 223 336.
[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

Start Machine Learning ×

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

[START MY EMAIL COURSE](#)