

[Click to Take the FREE Imbalanced Classification Crash-Course](#)



ROC Curves and Precision-Recall Curves for Imbalanced Classification

by Jason Brownlee on [January 6, 2020](#) in [Imbalanced Classification](#)

[Tweet](#)[Share](#)[Share](#)

Last Updated on January 14, 2020

Most imbalanced classification problems involve two classes: a negative case with the majority of examples and a positive case with a minority of examples.

Two diagnostic tools that help in the interpretation of binary (two-class) classification predictive models are ROC Curves and Precision-Recall curves.

Plots from the curves can be created and used to understand the trade-off in performance for different threshold values when interpreting probabilistic predictions. Each plot can also be summarized with an area under the curve score that can be used to directly compare classification models.

In this tutorial, you will discover ROC Curves and Precision-Recall Curves for imbalanced classification.

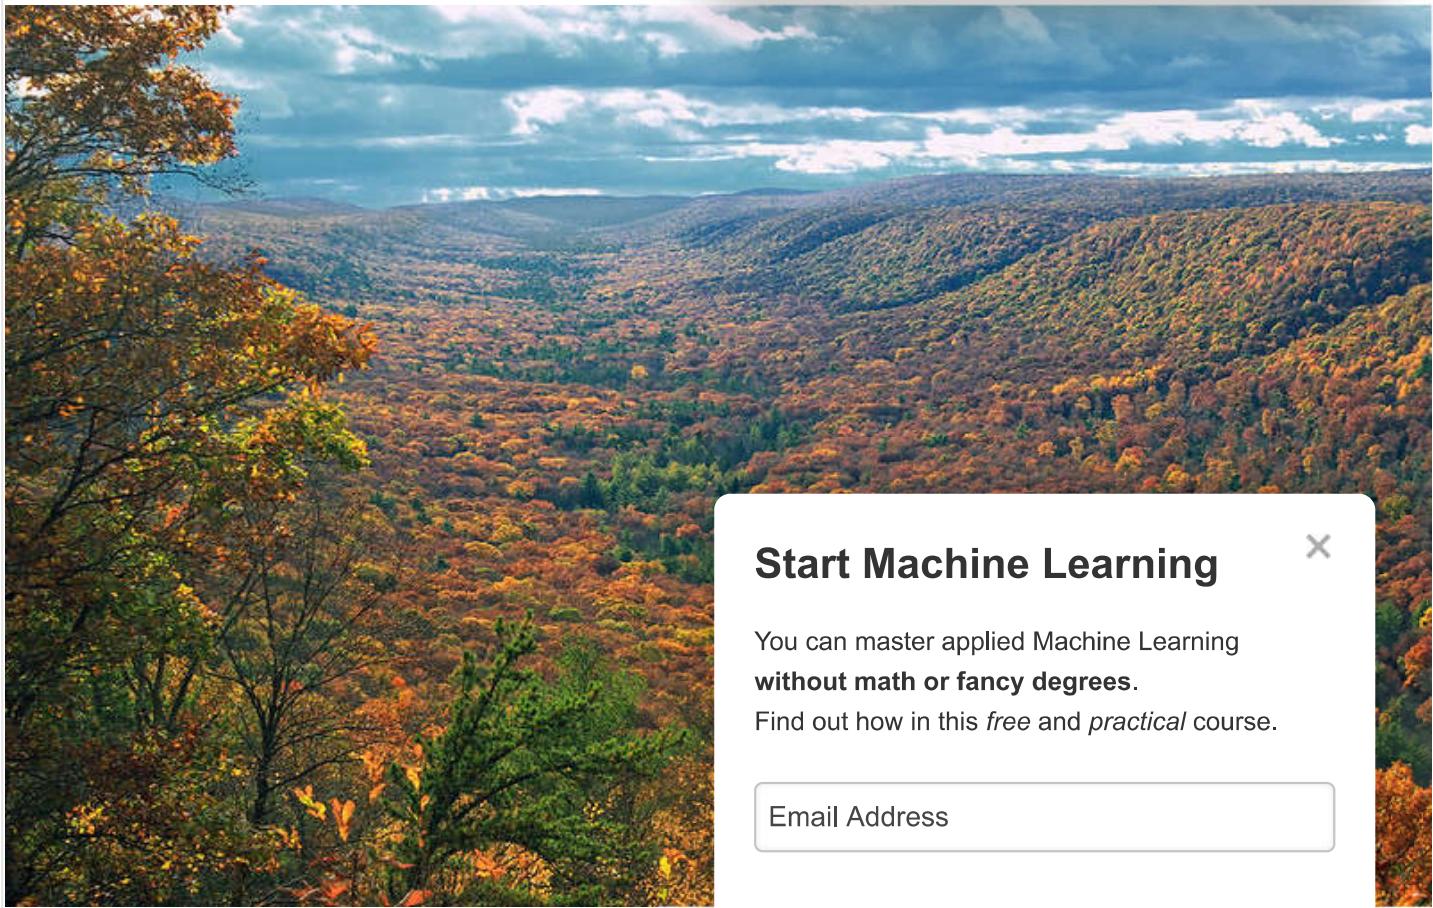
After completing this tutorial, you will know:

- ROC Curves and Precision-Recall Curves provide a diagnostic tool for binary classification models.
- ROC AUC and Precision-Recall AUC provide scores that summarize the curves and can be used to compare classifiers.
- ROC Curves and ROC AUC can be optimistic on severely imbalanced classification problems with few samples of the minority class.

Discover SMOTE, one-class classification, cost-sensitive learning, threshold moving, and much more [in my new book](#), with 30 step-by-step tutorials and full Python source code.

Let's get started.

[Start Machine Learning](#)



ROC Curves and Precision-Recall C

Photo by Nicholas A. Ton

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Tutorial Overview

This tutorial is divided into four parts; they are:

1. Review of the Confusion Matrix
2. ROC Curves and ROC AUC
3. Precision-Recall Curves and AUC
4. ROC and Precision-Recall Curves With a Severe Imbalance

Review of the Confusion Matrix

Before we dive into ROC Curves and PR Curves, it is important to review the confusion matrix.

For imbalanced classification problems, the majority class is typically referred to as the negative outcome (e.g. such as “*no change*” or “*negative test result*”), and the minority class is typically referred to as the positive outcome (e.g. “*change*” or “*positive test result*”).

The confusion matrix provides more insight into not only the performance of a predictive model, but also which classes are being predicted correctly, which incorrectly, and what type of errors are being made.

[Start Machine Learning](#)

The simplest confusion matrix is for a two-class classification problem, with negative (class 0) and positive (class 1) classes.

In this type of confusion matrix, each cell in the table has a specific and well-understood name, summarized as follows:

	Positive Prediction	Negative Prediction
Positive Class	True Positive (TP)	False Negative (FN)
Negative Class	False Positive (FP)	True Negative (TN)

The metrics that make up the ROC curve and the precision-recall curve are defined in terms of the cells in the confusion matrix.

Now that we have brushed up on the confusion matrix

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

ROC Curves and ROC AUC

An ROC curve (or receiver operating characteristic curve) is a plot that summarizes the performance of a binary classification model on the positive class.

The x-axis indicates the False Positive Rate and the y-axis indicates the True Positive Rate.

- **ROC Curve:** Plot of False Positive Rate (x) vs. True Positive Rate (y).

The true positive rate is a fraction calculated as the total number of true positive predictions divided by the sum of the true positives and the false negatives (e.g. all examples in the positive class). The true positive rate is referred to as the sensitivity or the recall.

- **TruePositiveRate** = $\text{TruePositives} / (\text{TruePositives} + \text{False Negatives})$

The false positive rate is calculated as the total number of false positive predictions divided by the sum of the false positives and true negatives (e.g. all examples in the negative class).

- **FalsePositiveRate** = $\text{FalsePositives} / (\text{FalsePositives} + \text{TrueNegatives})$

Start Machine Learning

We can think of the plot as the fraction of correct predictions for the positive class (y-axis) versus the fraction of errors for the negative class (x-axis).

Ideally, we want the fraction of correct positive class predictions to be 1 (top of the plot) and the fraction of incorrect negative class predictions to be 0 (left of the plot). This highlights that the best possible classifier that achieves perfect skill is the top-left of the plot (coordinate 0,1).

- **Perfect Skill:** A point in the top left of the plot.

The threshold is applied to the cut-off point in probability between the positive and negative classes, which by default for any classifier would be set at 0.5, halfway between each outcome (0 and 1).

A trade-off exists between the TruePositiveRate and FalsePositiveRate. The balance of predictions of classification will change the balance of predictions in expense of FalsePositiveRate, or the reverse case.

By evaluating the true positive and false positives for different threshold values, a curve that stretches from the bottom left to top right and bows inward is formed.

A classifier that has no discriminative power between two classes will have a diagonal line between a False Positive Rate of 0 and a True Positive Rate of 1 (from a False Negative Rate of 1 and a True Negative Rate of 0). Models represented by points below this line have some discriminative power.

The curve provides a convenient diagnostic tool to investigate one classifier with different threshold values and the effect on the TruePositiveRate and FalsePositiveRate. One might choose a threshold in order to bias the predictive behavior of a classification model.

It is a popular diagnostic tool for classifiers on balanced and imbalanced binary prediction problems alike because it is not biased to the majority or minority class.

 *ROC analysis does not have any bias toward models that perform well on the majority class at the expense of the majority class—a property that is quite attractive when dealing with imbalanced data.*

— Page 27, [Imbalanced Learning: Foundations, Algorithms, and Applications](#), 2013.

We can plot a ROC curve for a model in Python using the `roc_curve()` scikit-learn function.

The function takes both the true outcomes (0,1) from the test set and the predicted probabilities for the 1 class. The function returns the false positive rates for each threshold, true positive rates for each threshold and thresholds.

```
1 ...
2 # calculate roc curve
3 fpr, tpr, thresholds = roc_curve(testy, pos_pr)
```

[Start Machine Learning](#)

Most scikit-learn models can predict probabilities by calling the `predict_proba()` function.

This will return the probabilities for each class, for each sample in a test set, e.g. two numbers for each of the two classes in a binary classification problem. The probabilities for the positive class can be retrieved as the second column in this array of probabilities.

```
1 ...
2 # predict probabilities
3 yhat = model.predict_proba(testX)
4 # retrieve just the probabilities for the positive class
5 pos_probs = yhat[:, 1]
```

We can demonstrate this on a synthetic dataset and plot the ROC curve for a no skill classifier and a Logistic Regression model.

The `make_classification()` function can be used to create a synthetic dataset. It will create 1,000 examples for a binary classification problem. We will split the dataset into a train and test sets of equal size.

```
1 ...
2 # generate 2 class dataset
3 X, y = make_classification(n_samples=1000, n_classes=2)
4 # split into train/test sets
5 trainX, testX, trainy, testy = train_test_split(X, y)
```

A Logistic Regression model is a good model for demonstrating well-calibrated, as opposed to other machine learning models, in which case their probabilities may need to be calibrated first (e.g. an SVM).

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

```
1 ...
2 # fit a model
3 model = LogisticRegression(solver='lbfgs')
4 model.fit(trainX, trainy)
```

The complete example is listed below.

```
1 # example of a roc curve for a predictive model
2 from sklearn.datasets import make_classification
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.model_selection import train_test_split
5 from sklearn.metrics import roc_curve
6 from matplotlib import pyplot
7 # generate 2 class dataset
8 X, y = make_classification(n_samples=1000, n_classes=2, random_state=1)
9 # split into train/test sets
10 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
11 # fit a model
12 model = LogisticRegression(solver='lbfgs')
13 model.fit(trainX, trainy)
14 # predict probabilities
15 yhat = model.predict_proba(testX)
16 # retrieve just the probabilities for the positive class
17 pos_probs = yhat[:, 1]
18 # plot no skill roc curve
19 pyplot.plot([0, 1], [0, 1], linestyle='--', label='No Skill')
20 # calculate roc curve for model
21 fpr, tpr, _ = roc_curve(testy, pos_probs)
22 # plot model roc curve
```

Start Machine Learning

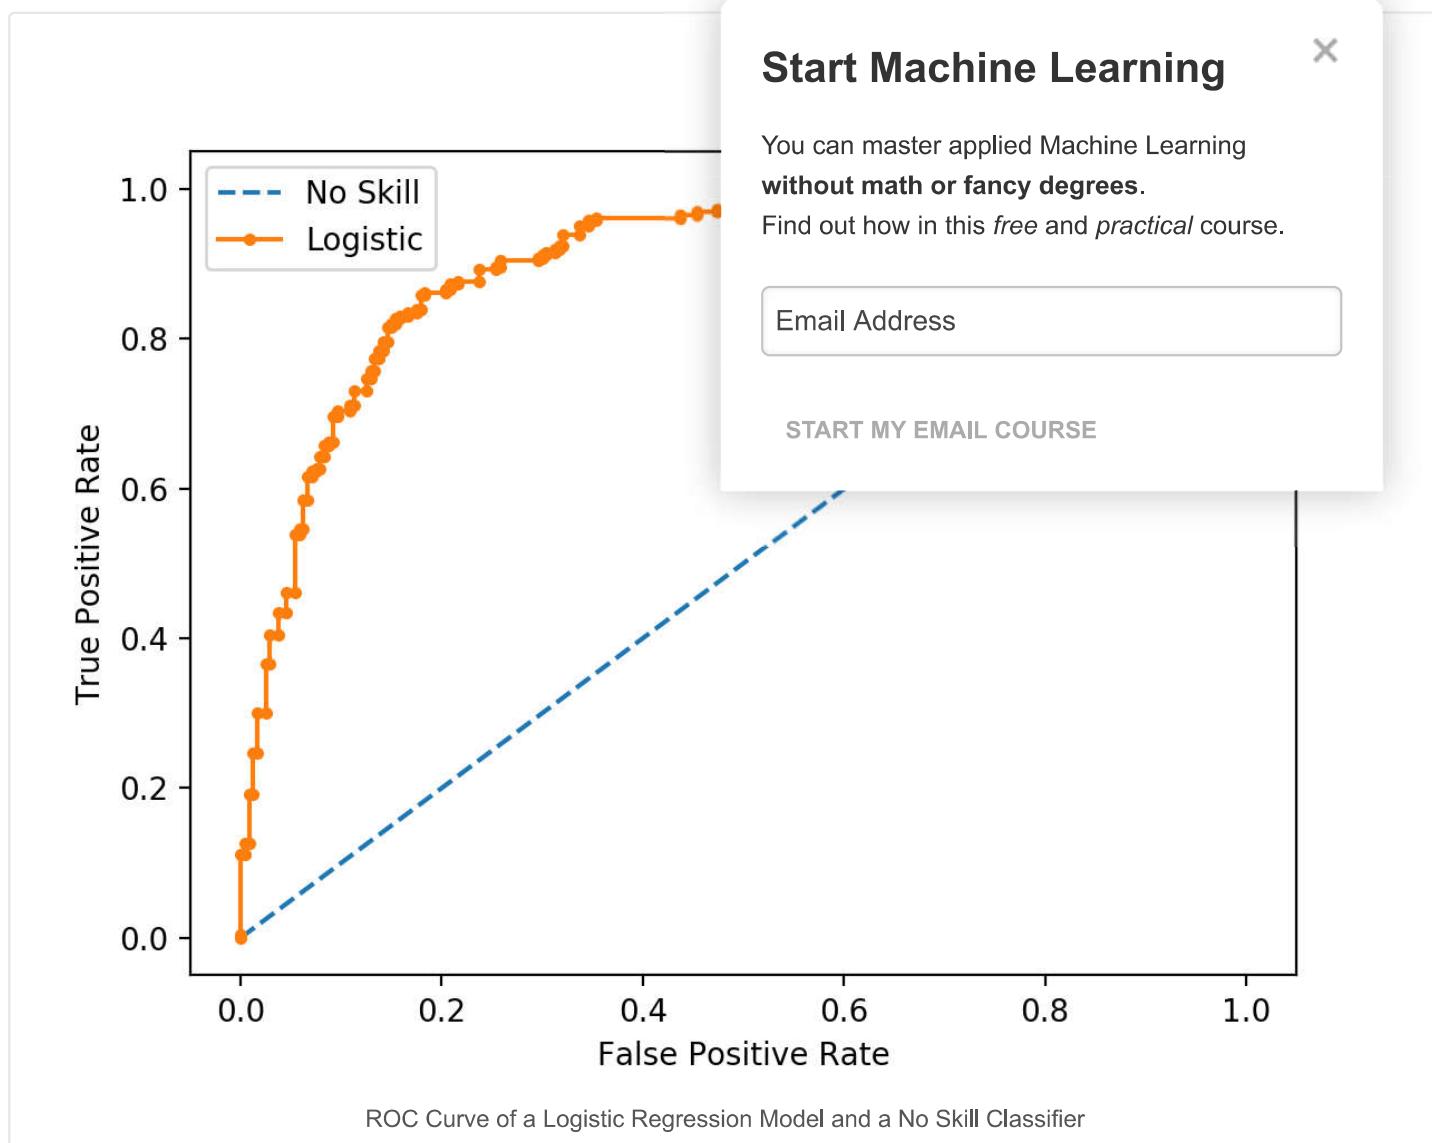
```

23 pyplot.plot(fpr, tpr, marker='.', label='Logistic')
24 # axis labels
25 pyplot.xlabel('False Positive Rate')
26 pyplot.ylabel('True Positive Rate')
27 # show the legend
28 pyplot.legend()
29 # show the plot
30 pyplot.show()

```

Running the example creates the synthetic dataset, splits into train and test sets, then fits a Logistic Regression model on the training dataset and uses it to make a prediction on the test set.

The ROC Curve for the Logistic Regression model is shown (orange with dots). A no skill classifier as a diagonal line (blue with dashes).



Now that we have seen the ROC Curve, let's take a closer look at the ROC area under curve score.

ROC Area Under Curve (AUC) Score

Although the ROC Curve is a helpful diagnostic tool, it can be challenging to compare two or more classifiers based on their curves.

[Start Machine Learning](#)

Instead, the area under the curve can be calculated to give a single score for a classifier model across all threshold values. This is called the ROC area under curve or ROC AUC or sometimes ROCAUC.

The score is a value between 0.0 and 1.0 for a perfect classifier.

AUCROC can be interpreted as the probability that the scores given by a classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one.

— Page 54, *Learning from Imbalanced Data Sets*, 2018.

This single score can be used to compare binary classifier models directly. As such, this score might be the most commonly used for comparing classification models.

The most common metric involves receiver operating characteristic (ROC) curves and the area under the ROC curve (AUC).

— Page 27, *Imbalanced Learning: Foundations, Algorithms, and Applications*.

The AUC for the ROC can be calculated in scikit-learn.

Like the `roc_curve()` function, the AUC function takes the true class labels and the predicted probabilities for the positive class.

```
1 ...
2 # calculate roc auc
3 roc_auc = roc_auc_score(testy, pos_probs)
```

We can demonstrate this the same synthetic dataset with a Logistic Regression model.

The complete example is listed below.

```
1 # example of a roc auc for a predictive model
2 from sklearn.datasets import make_classification
3 from sklearn.dummy import DummyClassifier
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import roc_auc_score
7 # generate 2 class dataset
8 X, y = make_classification(n_samples=1000, n_classes=2, random_state=1)
9 # split into train/test sets
10 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
11 # no skill model, stratified random class predictions
12 model = DummyClassifier(strategy='stratified')
13 model.fit(trainX, trainy)
14 yhat = model.predict_proba(testX)
15 pos_probs = yhat[:, 1]
16 # calculate roc auc
17 roc_auc = roc_auc_score(testy, pos_probs)
18 print('No Skill ROC AUC %.3f' % roc_auc)
19 # skilled model
20 model = LogisticRegression(solver='lbfgs')
21 model.fit(trainX, trainy)
22 yhat = model.predict_proba(testX)
```

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

the

Start Machine Learning

```
23 pos_probs = yhat[:, 1]
24 # calculate roc auc
25 roc_auc = roc_auc_score(testy, pos_probs)
26 print('Logistic ROC AUC %.3f' % roc_auc)
```

Running the example creates and splits the synthetic dataset, fits the model, and uses the fit model to predict probabilities on the test dataset.

In this case, we can see that the ROC AUC for the Logistic Regression model on the synthetic dataset is about 0.903, which is much better than a no skill classifier with a score of about 0.5.

```
1 No Skill ROC AUC 0.509
2 Logistic ROC AUC 0.903
```

Although widely used, the ROC AUC is not without problems.

For imbalanced classification with a severe skew and can be misleading. This is because a small number of change in the ROC Curve or ROC AUC score.

 *Although ROC graphs are widely used to evaluate imbalance, it has a drawback: under class rarity associated to the presence of a low sample size, ROC curves are unreliable.*

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

— Page 55, *Learning from Imbalanced Data Sets*, 2018.

A common alternative is the precision-recall curve and area under curve.

Precision-Recall Curves and AUC

Precision is a metric that quantifies the number of correct positive predictions made.

It is calculated as the number of true positives divided by the total number of true positives and false positives.

- **Precision** = TruePositives / (TruePositives + FalsePositives)

The result is a value between 0.0 for no precision and 1.0 for full or perfect precision.

Recall is a metric that quantifies the number of correct positive predictions made out of all positive predictions that could have been made.

It is calculated as the number of true positives divided by the total number of true positives and false negatives (e.g. it is the true positive rate).

- **Recall** = TruePositives / (TruePositives + FalseNegatives)

Start Machine Learning

The result is a value between 0.0 for no recall and 1.0 for full or perfect recall.

Both the precision and the recall are focused on the positive class (the minority class) and are unconcerned with the true negatives (majority class).

“ ... precision and recall make it possible to assess the performance of a classifier on the minority class.

— Page 27, **Imbalanced Learning: Foundations, Algorithms, and Applications**, 2013.

A precision-recall curve (or PR Curve) is a plot of the precision (y-axis) and the recall (x-axis) for different probability thresholds.

- **PR Curve:** Plot of Recall (x) vs Precision (y).

A model with perfect skill is depicted as a point at a coordinate of (1,1). A no-skew classifier has a precision that is proportional to the number of positive samples. The minimum precision will be 0.5.

The focus of the PR curve on the minority class makes it particularly useful for classification models.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

“ Precision-recall curves (PR curves) are recommended for highly skewed domains where ROC curves may provide an excessively optimistic view of the performance.

— **A Survey of Predictive Modelling under Imbalanced Distributions**, 2015.

A precision-recall curve can be calculated in scikit-learn using the `precision_recall_curve()` function that takes the class labels and predicted probabilities for the minority class and returns the precision, recall, and thresholds.

```
1 ...
2 # calculate precision-recall curve
3 precision, recall, _ = precision_recall_curve(testy, pos_probs)
```

We can demonstrate this on a synthetic dataset for a predictive model.

The complete example is listed below.

```
1 # example of a precision-recall curve for a predictive model
2 from sklearn.datasets import make_classification
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.model_selection import train_test_split
5 from sklearn.metrics import precision_recall_curve
6 from matplotlib import pyplot
7 # generate 2 class dataset
8 X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, n_clusters=2, weights=[0.9, 0.1], random_state=42)
9 # split into train/test sets
```

Start Machine Learning

```

10 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
11 # fit a model
12 model = LogisticRegression(solver='lbfgs')
13 model.fit(trainX, trainy)
14 # predict probabilities
15 yhat = model.predict_proba(testX)
16 # retrieve just the probabilities for the positive class
17 pos_probs = yhat[:, 1]
18 # calculate the no skill line as the proportion of the positive class
19 no_skill = len(y[y==1]) / len(y)
20 # plot the no skill precision-recall curve
21 pyplot.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
22 # calculate model precision-recall curve
23 precision, recall, _ = precision_recall_curve(testy, pos_probs)
24 # plot the model precision-recall curve
25 pyplot.plot(recall, precision, marker='.', label='Logistic')
26 # axis labels
27 pyplot.xlabel('Recall')
28 pyplot.ylabel('Precision')
29 # show the legend
30 pyplot.legend()
31 # show the plot
32 pyplot.show()

```

Running the example creates the synthetic dataset, splits it into training and testing sets, fits a Logistic Regression model on the training dataset and uses it to make predictions on the test dataset.

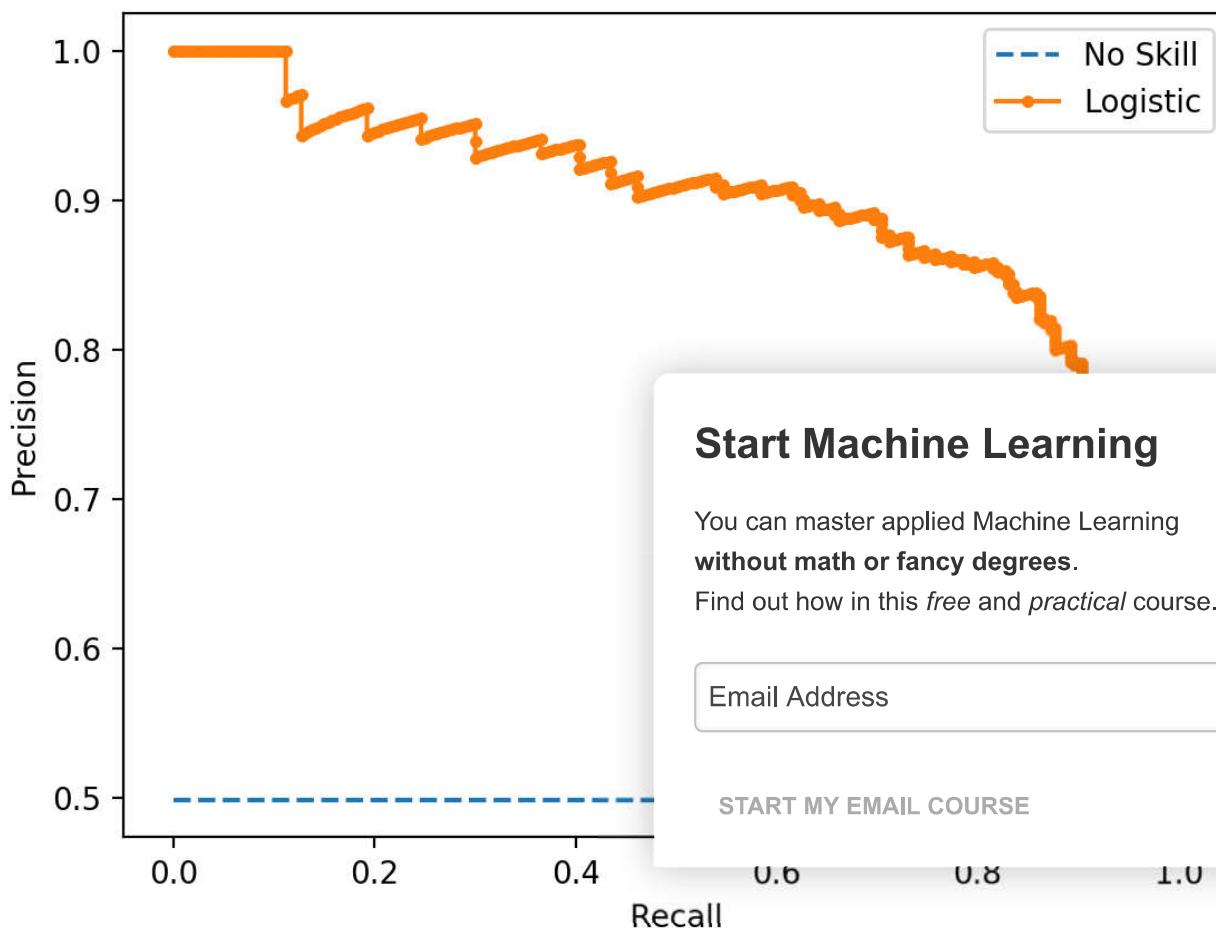
The Precision-Recall Curve for the Logistic Regression model is shown below. A horizontal blue line represents the baseline classifier is shown as a horizontal line (blue vertical line).

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Precision-Recall Curve of a Logistic Regression Model and a No Skill Classifier

Now that we have seen the Precision-Recall Curve, let's take a closer look at the ROC area under curve score.

Precision-Recall Area Under Curve (AUC) Score

The Precision-Recall AUC is just like the ROC AUC, in that it summarizes the curve with a range of threshold values as a single score.

The score can then be used as a point of comparison between different models on a binary classification problem where a score of 1.0 represents a model with perfect skill.

The Precision-Recall AUC score can be calculated using the `auc()` function in scikit-learn, taking the precision and recall values as arguments.

```

1 ...
2 # calculate the precision-recall auc
3 auc_score = auc(recall, precision)

```

Again, we can demonstrate calculating the Precision-Recall AUC score using the `precision_recall_curve()` function on the `liver-disease` dataset.

[Start Machine Learning](#)

The complete example is listed below.

```
1 # example of a precision-recall auc for a predictive model
2 from sklearn.datasets import make_classification
3 from sklearn.dummy import DummyClassifier
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import precision_recall_curve
7 from sklearn.metrics import auc
8 # generate 2 class dataset
9 X, y = make_classification(n_samples=1000, n_classes=2, random_state=1)
10 # split into train/test sets
11 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
12 # no skill model, stratified random class predictions
13 model = DummyClassifier(strategy='stratified')
14 model.fit(trainX, trainy)
15 yhat = model.predict_proba(testX)
16 pos_probs = yhat[:, 1]
17 # calculate the precision-recall auc
18 precision, recall, _ = precision_recall_curve
19 auc_score = auc(recall, precision)
20 print('No Skill PR AUC: %.3f' % auc_score)
21 # fit a model
22 model = LogisticRegression(solver='lbfgs')
23 model.fit(trainX, trainy)
24 yhat = model.predict_proba(testX)
25 pos_probs = yhat[:, 1]
26 # calculate the precision-recall auc
27 precision, recall, _ = precision_recall_curve
28 auc_score = auc(recall, precision)
29 print('Logistic PR AUC: %.3f' % auc_score)
```

Running the example creates and splits the synthetic dataset, fits the model, and uses the fit model to predict probabilities on the test dataset.

In this case, we can see that the Precision-Recall AUC for the Logistic Regression model on the synthetic dataset is about 0.898, which is much better than a no skill classifier that would achieve the score in this case of 0.632.

```
1 No Skill PR AUC: 0.632
2 Logistic PR AUC: 0.898
```

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

ROC and Precision-Recall Curves With a Severe Imbalance

In this section, we will explore the case of using the ROC Curves and Precision-Recall curves with a binary classification problem that has a severe class imbalance.

Firstly, we can use the `make_classification()` function to create 1,000 examples for a classification problem with about a 1:100 minority to majority class ratio. This can be achieved by setting the “`weights`” argument and specifying the weighting of generated instances from each class.

We will use a 99 percent and 1 percent weighting with 1,000 total examples, meaning there would be about 990 for class 0 and about 10 for class 1.

```
1 ...
2 # generate 2 class dataset
```

Start Machine Learning

```
3 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], random_state=1)
```

We can then split the dataset into training and test sets and ensure that both have the same general class ratio by setting the “*stratify*” argument on the call to the *train_test_split()* function and setting it to the array of target variables.

```
1 ...
2 # split into train/test sets with same class ratio
3 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
```

Tying this together, the complete example of preparing the imbalanced dataset is listed below.

```
1 # create an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import train_test_split
4 # generate 2 class dataset
5 X, y = make_classification(n_samples=1000, n_features=2, n_classes=2, weights=[0.99, 0.01], random_state=1)
6 # split into train/test sets with same class ratio
7 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
8 # summarize dataset
9 print('Dataset: Class0=%d, Class1=%d' % (len(trainy), len(testy)))
10 print('Train: Class0=%d, Class1=%d' % (len(trainX), len(trainy)))
11 print('Test: Class0=%d, Class1=%d' % (len(testX), len(testy)))
```

Running the example first summarizes the class ratio for the dataset, then splits it into three parts: train and test sets, confirming the split of the dataset has been performed correctly.

```
1 Dataset: Class0=985, Class1=15
2 Train: Class0=492, Class1=8
3 Test: Class0=493, Class1=7
```

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this free and practical course.

START MY EMAIL COURSE

Next, we can develop a Logistic Regression model on the dataset and evaluate the performance of the model using a ROC Curve and ROC AUC score, and compare the results to a no skill classifier, as we did in a prior section.

The complete example is listed below.

```
1 # roc curve and roc auc on an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.dummy import DummyClassifier
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import roc_curve
7 from sklearn.metrics import roc_auc_score
8 from matplotlib import pyplot
9
10 # plot no skill and model roc curves
11 def plot_roc_curve(test_y, naive_probs, model_probs):
12     # plot naive skill roc curve
13     fpr, tpr, _ = roc_curve(test_y, naive_probs)
14     pyplot.plot(fpr, tpr, linestyle='--', label='No Skill')
15     # plot model roc curve
16     fpr, tpr, _ = roc_curve(test_y, model_probs)
17     pyplot.plot(fpr, tpr, marker='.', label='Logistic')
18     # axis labels
19     pyplot.xlabel('False Positive Rate')
20     pyplot.ylabel('True Positive Rate')
21     # show the legend
22     pyplot.legend()
```

Start Machine Learning

```

23     # show the plot
24     pyplot.show()
25
26 # generate 2 class dataset
27 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], random_state=1)
28 # split into train/test sets with same class ratio
29 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
30 # no skill model, stratified random class predictions
31 model = DummyClassifier(strategy='stratified')
32 model.fit(trainX, trainy)
33 yhat = model.predict_proba(testX)
34 naive_probs = yhat[:, 1]
35 # calculate roc auc
36 roc_auc = roc_auc_score(testy, naive_probs)
37 print('No Skill ROC AUC %.3f' % roc_auc)
38 # skilled model
39 model = LogisticRegression(solver='lbfgs')
40 model.fit(trainX, trainy)
41 yhat = model.predict_proba(testX)
42 model_probs = yhat[:, 1]
43 # calculate roc auc
44 roc_auc = roc_auc_score(testy, model_probs)
45 print('Logistic ROC AUC %.3f' % roc_auc)
46 # plot roc curves
47 plot_roc_curve(testy, naive_probs, model_probs)

```

Running the example creates the imbalanced binary classification problem.

Then a logistic regression model is fit on the training data and the ROC curve for the model and the no skill classifier is evaluated alongside for reference.

The ROC AUC scores for both classifiers are reported, showing the no skill classifier achieving the lowest score of approximately 0.5 as expected. The results for the logistic regression model suggest it has some skill with a score of about 0.869.

1	No Skill ROC AUC 0.490
2	Logistic ROC AUC 0.869

A ROC curve is also created for the model and the no skill classifier, showing not excellent performance, but definitely skillful performance as compared to the diagonal no skill.

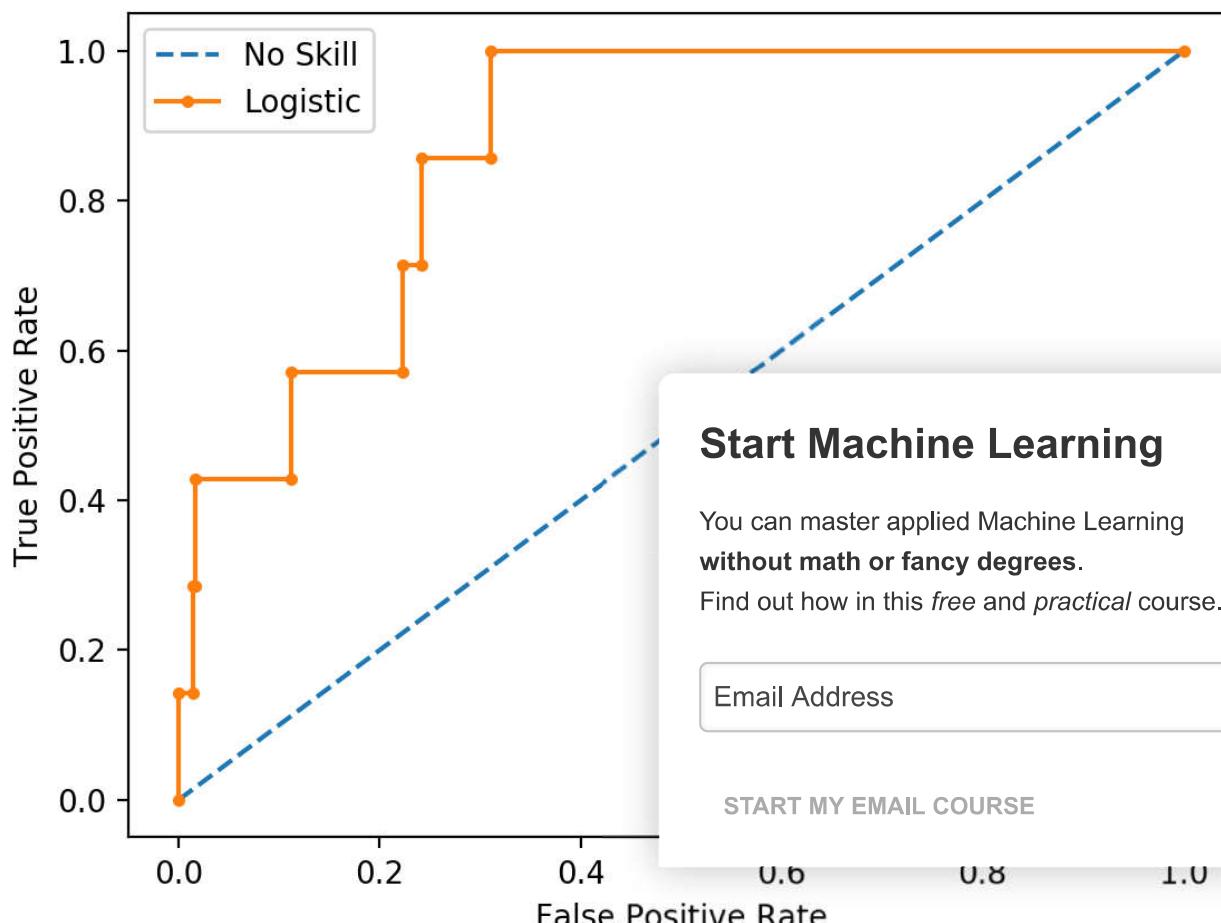
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)



Plot of ROC Curve for Logistic Regression on Imbalanced Classification Dataset

Next, we can perform an analysis of the same model fit and evaluated on the same data using the precision-recall curve and AUC score.

The complete example is listed below.

```

1 # pr curve and pr auc on an imbalanced dataset
2 from sklearn.datasets import make_classification
3 from sklearn.dummy import DummyClassifier
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import precision_recall_curve
7 from sklearn.metrics import auc
8 from matplotlib import pyplot
9
10 # plot no skill and model precision-recall curves
11 def plot_pr_curve(test_y, model_probs):
12     # calculate the no skill line as the proportion of the positive class
13     no_skill = len(test_y[test_y==1]) / len(test_y)
14     # plot the no skill precision-recall curve
15     pyplot.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
16     # plot model precision-recall curve
17     precision, recall, _ = precision_recall_curve(test_y, model_probs)
18     pyplot.plot(recall, precision, marker='.')
19     # axis labels

```

Start Machine Learning

```

20     pyplot.xlabel('Recall')
21     pyplot.ylabel('Precision')
22     # show the legend
23     pyplot.legend()
24     # show the plot
25     pyplot.show()
26
27 # generate 2 class dataset
28 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], random_state=1)
29 # split into train/test sets with same class ratio
30 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
31 # no skill model, stratified random class predictions
32 model = DummyClassifier(strategy='stratified')
33 model.fit(trainX, trainy)
34 yhat = model.predict_proba(testX)
35 naive_probs = yhat[:, 1]
36 # calculate the precision-recall auc
37 precision, recall, _ = precision_recall_curve(ytest, naive_probs)
38 auc_score = auc(recall, precision)
39 print('No Skill PR AUC: %.3f' % auc_score)
40 # fit a model
41 model = LogisticRegression(solver='lbfgs')
42 model.fit(trainX, trainy)
43 yhat = model.predict_proba(testX)
44 model_probs = yhat[:, 1]
45 # calculate the precision-recall auc
46 precision, recall, _ = precision_recall_curve(ytest, model_probs)
47 auc_score = auc(recall, precision)
48 print('Logistic PR AUC: %.3f' % auc_score)
49 # plot precision-recall curves
50 plot_pr_curve(testy, model_probs)

```

As before, running the example creates the imbalance.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

In this case we can see that the Logistic Regression model achieves a PR AUC of about 0.228 and a no skill model achieves a PR AUC of about 0.007.

```

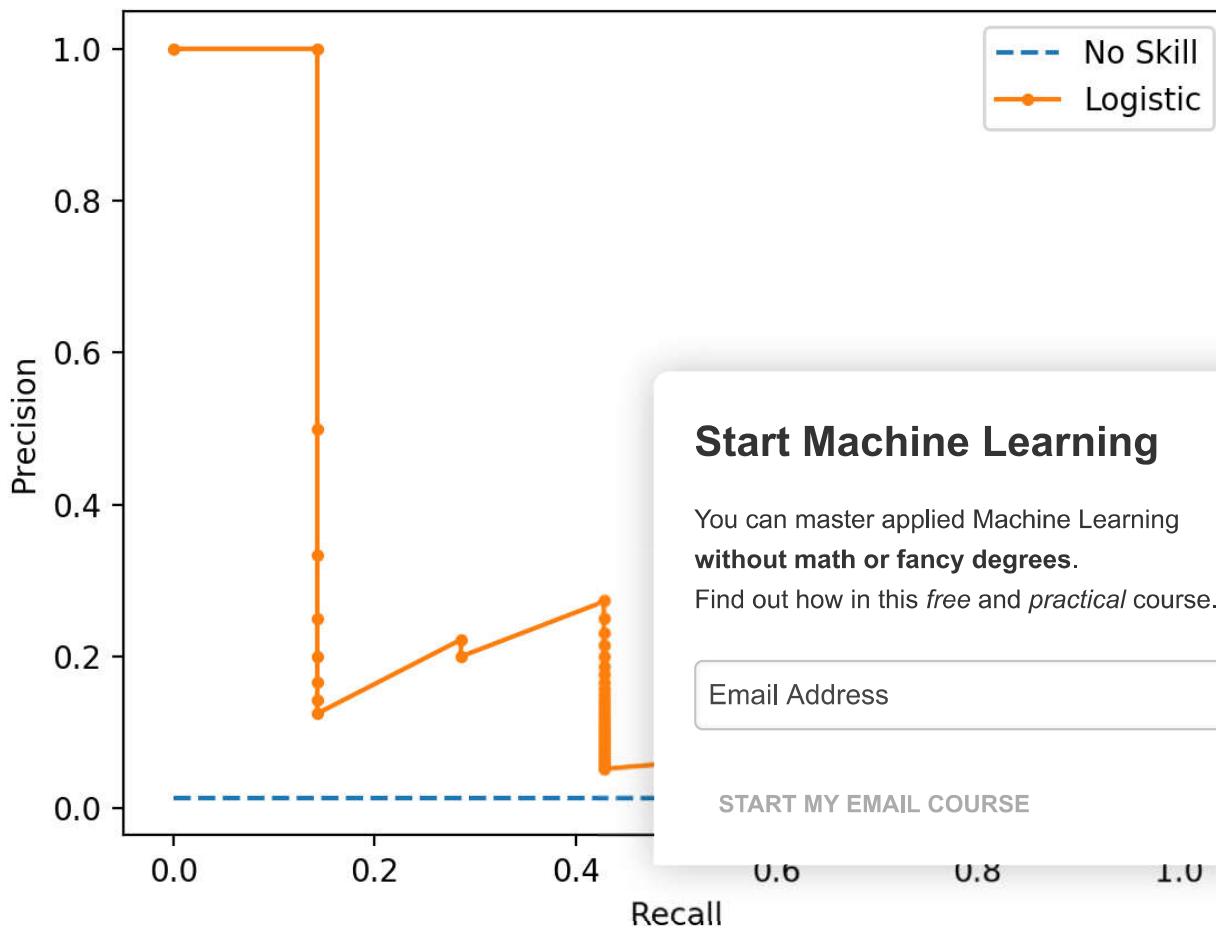
1 No Skill PR AUC: 0.007
2 Logistic PR AUC: 0.228

```

A plot of the precision-recall curve is also created.

We can see the horizontal line of the no skill classifier as expected and in this case the zig-zag line of the logistic regression curve close to the no skill line.

[Start Machine Learning](#)



To explain why the ROC and PR curves tell a different story, recall that the PR curve focuses on the minority class, whereas the ROC curve covers both classes.

If we use a threshold of 0.5 and use the logistic regression model to make a prediction for all examples in the test set, we see that it predicts class 0 or the majority class in all cases. This can be confirmed by using the `fit` model to predict crisp class labels, that will use the default threshold of 0.5. The distribution of predicted class labels can then be summarized.

```

1 ...
2 # predict class labels
3 yhat = model.predict(testX)
4 # summarize the distribution of class labels
5 print(Counter(yhat))

```

We can then create a histogram of the predicted probabilities of the positive class to confirm that the mass of predicted probabilities is below 0.5, and therefore are mapped to class 0.

```

1 ...
2 # create a histogram of the predicted probabilities
3 pyplot.hist(pos_probs, bins=100)
4 pyplot.show()

```

Start Machine Learning

Tying this together, the complete example is listed below.

```
1 # summarize the distribution of predicted probabilities
2 from collections import Counter
3 from matplotlib import pyplot
4 from sklearn.datasets import make_classification
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.model_selection import train_test_split
7 # generate 2 class dataset
8 X, y = make_classification(n_samples=1000, n_classes=2, weights=[0.99, 0.01], random_state=1)
9 # split into train/test sets with same class ratio
10 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
11 # fit a model
12 model = LogisticRegression(solver='lbfgs')
13 model.fit(trainX, trainy)
14 # predict probabilities
15 yhat = model.predict_proba(testX)
16 # retrieve just the probabilities for the positive class
17 pos_probs = yhat[:, 1]
18 # predict class labels
19 yhat = model.predict(testX)
20 # summarize the distribution of class labels
21 print(Counter(yhat))
22 # create a histogram of the predicted probabilities
23 pyplot.hist(pos_probs, bins=100)
24 pyplot.show()
```

Running the example first summarizes the distribution of predicted class labels. The majority class (class 0) is predicted for all examples in

```
1 Counter({0: 500})
```

A histogram plot of the predicted probabilities for class 1 is also created, showing the center of mass (most predicted probabilities) is less than 0.5 and in fact is generally close to zero.

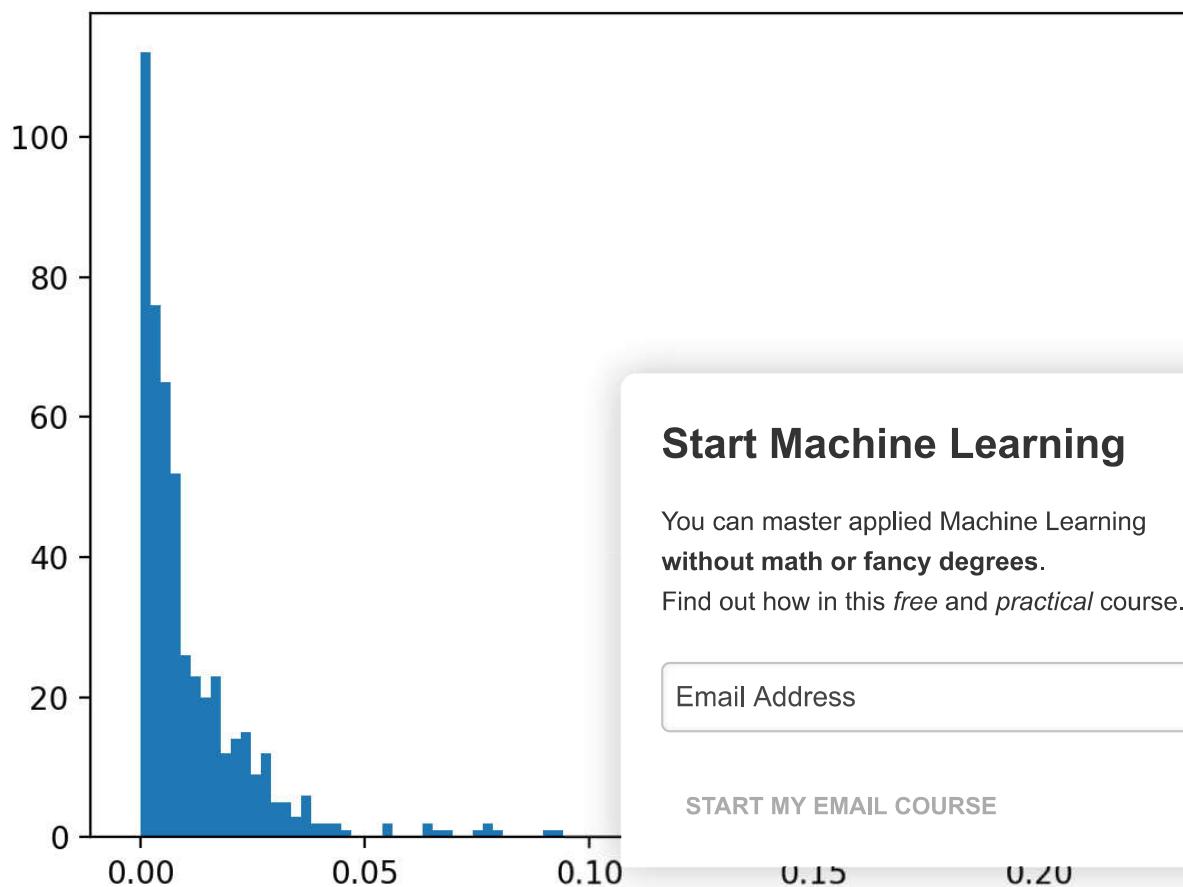
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning



Histogram of Logistic Regression Predicted Probabilities for Class 1 for Imbalanced Classification

Start Machine Learning ×

You can master applied Machine Learning **without math or fancy degrees**.
Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

This means, unless probability threshold is carefully chosen, any skillful nuance in the predictions made by the model will be lost. Selecting thresholds used to interpret predicted probabilities as crisp class labels is an important topic

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Tutorials

- [How to Use ROC Curves and Precision-Recall Curves for Classification in Python](#)

Papers

- [A Survey of Predictive Modelling under Imbalanced Distributions, 2015.](#)

Books

- [Imbalanced Learning: Foundations, Algorithms, and Applications](#)

Start Machine Learning

- Learning from Imbalanced Data Sets, 2018.

API

- `sklearn.datasets.make_classification` API.
- `sklearn.metrics.roc_curve` API.
- `sklearn.metrics.roc_auc_score` API
- `precision_recall_curve` API.
- `sklearn.metrics.auc` API.

Articles

- Receiver operating characteristic, Wikipedia.
- Precision and recall, Wikipedia.

Summary

In this tutorial, you discovered ROC Curves and Precision-Recall Curves.

Specifically, you learned:

- ROC Curves and Precision-Recall Curves provide a visual way to compare classifiers.
- ROC AUC and Precision-Recall AUC provide scores that can be used to compare classifiers.
- ROC Curves and ROC AUC can be optimistic on severely imbalanced classification problems with few samples of the minority class.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Get a Handle on Imbalanced Classification!

Develop Imbalanced Learning Models in Minutes

...with just a few lines of python code

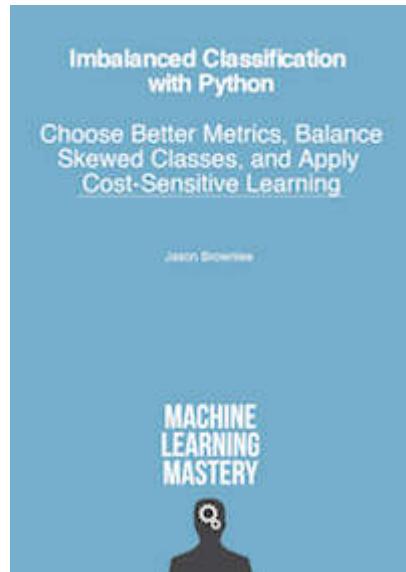
Discover how in my new Ebook:
[Imbalanced Classification with Python](#)

It provides **self-study tutorials** and **end-to-end projects** on:

Performance Metrics, Undersampling Methods, SMOTE, Threshold Moving, Probability Calibration, Cost-Sensitive Algorithms
and much more...

Bring Imbalanced Classification Methods to Your Machine Learning Projects

[Start Machine Learning](#)



SEE WHAT'S INSIDE

Choose Better Metrics, Balance Skewed Classes, and Apply Cost-Sensitive Learning

Jason Brownlee

MACHINE
LEARNING
MASTERY



Tweet

Share

Share



About Jason Brownlee

Jason Brownlee, PhD is a machine learning researcher and author of the book *Machine Learning Mastery*. He has written several articles and books on modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

◀ How to Calculate Precision, Recall, and F-Measure for Imbalanced Classification

Tour of Evaluation Metrics for Imbalanced Classification >

20 Responses to *ROC Curves and Precision-Recall Curves for Imbalanced Classification*



Xingsheng Qian January 6, 2020 at 4:30 pm #

REPLY ↗

Thanks for sharing, it's a good reading!



Jason Brownlee January 7, 2020 at 7:17 am #

REPLY ↗

You're welcome!

Start Machine Learning



Temitope Mamukuyomi January 10, 2020 at 8:36 am #

REPLY ↗

Thanks for the tutorial.

1. Other metrics to check for model validation are the summary of the model containing the estimates.
What inference can we deduct from those estimates?
2. Reason why I said this is because if it were to be a real life situation and not random number generated from the machine.
3. Then, after discovering that the fitted model has a very low precision (i.e. <0.5) is it fair to go ahead to use this model for unseen data (predictions)?
If not, what is the way forward?



Jason Brownlee January 10, 2020 at 1:30 pm #

If a model performs better than a naive baseline, does that mean it's good?

You can decide whether you want to use it or continue to work on it to make it better.

Testing new models ends when you run out of time or project stakeholders.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Greg January 7, 2020 at 7:57 am #

REPLY ↗

Great article as usual! I appreciate all you do for the community.



Jason Brownlee January 7, 2020 at 1:45 pm #

REPLY ↗

Thanks Greg!



marco January 7, 2020 at 10:52 pm #

REPLY ↗

Hello Jason,

a question about Feature Extraction and Feature Selection.

Can they be classified as Unsupervised Machine Learning?

Can be useful using Feature Extraction and Feature Selection with Keras?

Can you please explain in easy words what are difference between both and what they are for?

Do you have any simple example?

Thanks a lot,

Marco

Start Machine Learning



Jason Brownlee January 8, 2020 at 8:26 am #

REPLY ↩

Maybe. They are not really a learning algorithm.

Yes, but it depends on the type of problem.

Feature selection controls which inputs from the data to feed into the model. Feature extraction provides a view or projection of the raw features as input to the model. Most neural nets will perform feature extraction automatically, e.g. a CNN for image classification.



marco January 8, 2020 at 6:40 pm #

X

Thanks Jason.
as far I can understand it seems to me that there is some
extraction. Right?
and the result in the extraction phase maybe the same
i.e. the dataset in the extraction phase maybe the same
Thanks

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

REPLY ↩



Jason Brownlee January 9, 2020 at 7:24 am #

Perhaps. It really depends on the data and the models.

For example, ensembles of decision trees perform an automatic type of feature selection. Neural nets perform an automatic type of feature extraction, etc.

With audio/txt data and standard ml algorithms, feature extraction is pretty much a required step. On text data, not at all, unless you call it feature engineering.

So on. It's complicated.



Markus January 9, 2020 at 8:08 am #

REPLY ↩

Hi

What do you exactly mean with "crisp class labels"? Especially the word crisp?

Thanks

Jason Brownlee January 9, 2020 at 8:15 am #

Start Machine Learning



Good question.

I mean label per sample, as opposed to probability of class membership for each sample.



Mathieu January 10, 2020 at 8:16 am #

REPLY ↗

Hi,

Please, it will be great if you can show us how to optimize the best threshold (lower than 0.5) for minimize False negative and False positive.



Jason Brownlee January 10, 2020

I have a tutorial on exactly this wr



Jay February 20, 2020 at 1:30 pm #

Do you have a special offer/discount code for

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee February 21, 2020 at 8:16 am #

REPLY ↗

Yes, see this:

<https://machinelearningmastery.com/faq/single-faq/can-i-have-a-discount>



Beco March 20, 2020 at 8:31 pm #

REPLY ↗

How would we interpret a case In which we get perfect accuracy but zero roc-auc, f1-score, precision and recall?



Jason Brownlee March 21, 2020 at 8:20 am #

REPLY ↗

It does not make sense. You may have a bug somewhere.



Jon April 1, 2020 at 4:36 am #

REPLY ↗

Start Machine Learning

Hello,

Many thanks for the great article which I found super useful.

I have a question. When you stated in the very last sentence of the article "This means, unless probability threshold is carefully chosen, any skillful nuance in the predictions made by the model will be lost. Selecting thresholds used to interpret predicted probabilities as crisp class labels is an important topic", therefore in this case, a lower probability than 0.5 needs to be taken right? For example 0.2.

Does this makes sense?

Do you have suggestions on how to pick up the best threshold in this case, or is it a matter of sorting by class 1 probabilities and taking the top N items?

Thanks and looking forward for your reply.

Jon



Jason Brownlee April 1, 2020 at 5:54 am #

Yes, see this:

<https://machinelearningmastery.com/threshold-modification/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Leave a Reply

Name (required)

Email (will not be published) (required)

Website

SUBMIT COMMENT

Start Machine Learning



Welcome!

My name is Jason Brownlee PhD, and I help developers get results with machine learning.

[Read more](#)

Never miss a tutorial:



Picked for you:



[8 Tactics to Combat Imbalanced Classes in Your Machine Learning Model](#)



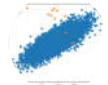
[SMOTE for Imbalanced Classification with Python](#)



[Imbalanced Classification With Python \(7-Day Mini-Course\)](#)



[Tour of Evaluation Metrics for Imbalanced Classification](#)



[One-Class Classification Algorithms for Imbalanced Datasets](#)

Start Machine Learning



You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

[START MY EMAIL COURSE](#)

© 2019 Machine Learning Mastery Pty. Ltd. All Rights Reserved.

Address: PO Box 206, Vermont Victoria 3133, Australia. | ACN: 626 223 336.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

[Start Machine Learning](#)

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Start Machine Learning